

```
In [12]: #EXP_8
```

```
In [13]: #Aim:Logistic Regression
```

```
In [1]: #Name:Dev Sanjay Vaidya  
#Roll no:69  
#Sec:B  
#Subject:ET-1  
#Date:11/09/25
```

To Perform And Analysis Of Logistic Regression Algorithm

Importing The Libraries¶

```
In [1]: import pandas as pd  
import numpy as np
```

Data Acquisitionuing Pandas

```
In [2]: import os
```

```
In [3]: os.getcwd()
```

```
Out[3]: 'C:\\\\Users\\\\LENOVO'
```

```
In [4]: os.chdir('C:\\\\Users\\\\LENOVO\\\\Desktop')
```

```
In [5]: data=pd.read_csv("heart.csv")
```

```
In [6]: data.head()
```

```
Out[6]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  
          0    52    1    0      125   212    0        1     168      0      1.0      2  
          1    53    1    0      140   203    1        0     155      1      3.1      0  
          2    70    1    0      145   174    0        1     125      1      2.6      0  
          3    61    1    0      148   203    0        1     161      0      0.0      2  
          4    62    0    0      138   294    1        1     106      0      1.9      1
```

```
In [7]: data.tail()
```

Out[7]:

| | age | sex | cp | trestbps | chol | fbس | restecg | thalach | exang | oldpeak | sl |
|-------------|-----|-----|----|----------|------|-----|---------|---------|-------|---------|----|
| 1020 | 59 | 1 | 1 | 140 | 221 | 0 | 1 | 164 | 1 | 0.0 | |
| 1021 | 60 | 1 | 0 | 125 | 258 | 0 | 0 | 141 | 1 | 2.8 | |
| 1022 | 47 | 1 | 0 | 110 | 275 | 0 | 0 | 118 | 1 | 1.0 | |
| 1023 | 50 | 0 | 0 | 110 | 254 | 0 | 0 | 159 | 0 | 0.0 | |
| 1024 | 54 | 1 | 0 | 120 | 188 | 0 | 1 | 113 | 0 | 1.4 | |

In [8]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         1025 non-null    int64  
 1   sex         1025 non-null    int64  
 2   cp          1025 non-null    int64  
 3   trestbps    1025 non-null    int64  
 4   chol        1025 non-null    int64  
 5   fbs         1025 non-null    int64  
 6   restecg     1025 non-null    int64  
 7   thalach     1025 non-null    int64  
 8   exang       1025 non-null    int64  
 9   oldpeak     1025 non-null    float64 
 10  slope       1025 non-null    int64  
 11  ca          1025 non-null    int64  
 12  thal        1025 non-null    int64  
 13  target      1025 non-null    int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

In [9]: `data.describe()`

Out[9]:

| | age | sex | cp | trestbps | chol |
|--------------|-------------|-------------|-------------|-------------|-------------|
| count | 1025.000000 | 1025.000000 | 1025.000000 | 1025.000000 | 1025.000000 |
| mean | 54.434146 | 0.695610 | 0.942439 | 131.611707 | 246.000000 |
| std | 9.072290 | 0.460373 | 1.029641 | 17.516718 | 51.59251 |
| min | 29.000000 | 0.000000 | 0.000000 | 94.000000 | 126.000000 |
| 25% | 48.000000 | 0.000000 | 0.000000 | 120.000000 | 211.000000 |
| 50% | 56.000000 | 1.000000 | 1.000000 | 130.000000 | 240.000000 |
| 75% | 61.000000 | 1.000000 | 2.000000 | 140.000000 | 275.000000 |
| max | 77.000000 | 1.000000 | 3.000000 | 200.000000 | 564.000000 |

In [10]: `data.shape`

```
Out[10]: (1025, 14)
```

```
In [11]: data.size
```

```
Out[11]: 14350
```

```
In [12]: data.ndim
```

```
Out[12]: 2
```

Data preprocessing _ data cleaning _ missing value treatment

```
In [13]: # check Missing Value by record
```

```
data.isna()
```

```
Out[13]:
```

| | age | sex | cp | trestbps | chol | fbs | restecg | thalach | exang | oldpeak |
|-------------|-------|-------|-------|----------|-------|-------|---------|---------|-------|---------|
| 0 | False | False | False | False | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1020 | False | False | False | False | False | False | False | False | False | False |
| 1021 | False | False | False | False | False | False | False | False | False | False |
| 1022 | False | False | False | False | False | False | False | False | False | False |
| 1023 | False | False | False | False | False | False | False | False | False | False |
| 1024 | False | False | False | False | False | False | False | False | False | False |

1025 rows × 14 columns

```
In [14]: data.isna().any()
```

```
Out[14]: age      False  
          sex      False  
          cp       False  
          trestbps False  
          chol      False  
          fbs       False  
          restecg   False  
          thalach   False  
          exang     False  
          oldpeak   False  
          slope     False  
          ca        False  
          thal      False  
          target    False  
          dtype: bool
```

```
In [15]: data.isna().sum()
```

```
Out[15]: age      0  
          sex      0  
          cp       0  
          trestbps 0  
          chol      0  
          fbs       0  
          restecg   0  
          thalach   0  
          exang     0  
          oldpeak   0  
          slope     0  
          ca        0  
          thal      0  
          target    0  
          dtype: int64
```

Independent and Dependent Variables

```
In [16]: x=data.drop("target", axis=1)  
y=data["target"]
```

Splitting of DataSet into train and Test

```
In [17]: #splitting the data into training and testing data sets  
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2 ,random_st
```

Logistic Regression

```
In [18]: import warnings  
from sklearn.exceptions import ConvergenceWarning
```

```
In [19]: from sklearn.linear_model import LogisticRegression
```

```
In [20]: log = LogisticRegression()  
log.fit(x_train, y_train)
```

```
C:\Users\LENOVO\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.p  
y:465: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:  
https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regre  
ssion
```

```
n_iter_i = _check_optimize_result(
```

```
Out[20]: ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression()
```

```
In [21]: y_pred1 = log.predict(x_test)
```

```
In [22]: from sklearn.metrics import accuracy_score
```

```
In [23]: accuracy_score(y_test,y_pred1)
```


```
Out[23]: 0.7853658536585366
```

```
In [24]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.metrics import confusion_matrix
```

```
In [25]: cm = confusion_matrix(y_test, y_pred1)
```

```
In [26]: labels = np.unique(y_test) # Get unique class labels  
cm_df = pd.DataFrame(cm, index=labels, columns=labels)
```

```
In [27]: # Plot confusion matrix using seaborn  
plt.figure(figsize=(6, 4))  
sns.heatmap(cm_df, annot=True, fmt='d', cmap='Blues', linewidths=1, linecolor  
e='black')  
  
plt.xlabel("Predicted Label")  
plt.ylabel("True Label")  
plt.title("Confusion Matrix")  
plt.show()
```



In []: